

P1: Part-of-Speech (PoS) Tagging

Natural Language Understanding

Interuniversity Master's Degree in Artificial Intelligence
Academic Year 2025-2026

Part-of-speech tagging

Tagging each word with its grammatical category (noun, verb, adjective, etc.).

El	volcán	emitió	mucha	lava	durante	la	erupción
DET	NOUN	VERB	DET	NOUN	PREP	DEP	NOUN

This is the simplest approach to structured prediction, that is, to convert unstructured text into a structured representation that can be used to automatically extract information.


Part-of-speech tagging

First	ADV
come	VERB
,	PUNCT
first	ADV
serve	VERB
.	PUNCT

The idea is to preprocess the text with this information, obtaining useful input data for higher-level NLP tasks:

- Machine translation...
- Question answering...
- It is also valuable for interdisciplinary studies between computer science and linguistics, providing insights into how humans express themselves differently, for instance, based on the domain or the audience they are addressing.

Part-of-speech tagging



Why do we eat toast for breakfast, and then toast to good health at dinner?

What does the turkey we eat on Thanksgiving have to do with the country on the eastern Mediterranean?

Can you figure out how much your dinner will cost by counting the words on the menu?

In *The Language of Food*, Stanford University professor and MacArthur Fellow Dan Jurafsky peels away the mysteries from the foods we think we know.

Objectives

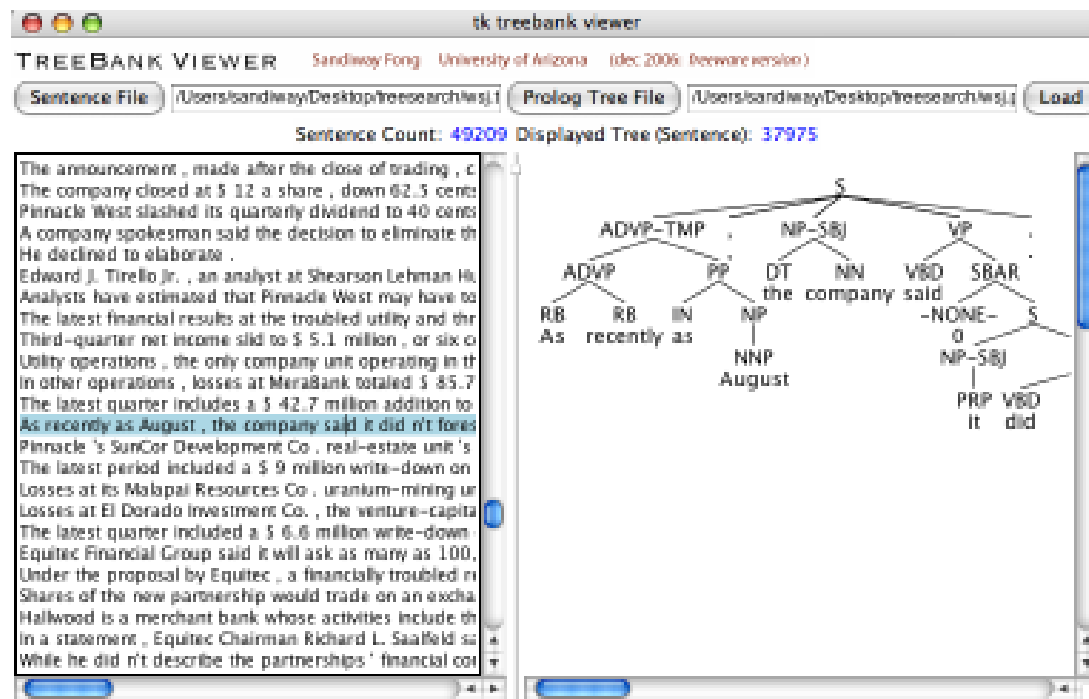
Work and process successfully Universal Dependency (UD) datasets

Build and train your own Part-of-Speech (PoS) taggers based on neural models

Compare neural architectures and their performance across natural languages

Treebanks

- In linguistics, a treebank is a parsed text corpus that annotates syntactic or semantic sentence structure



Mitchell P. Marcus, Beatrice Santorini, and Mary Ann Marcinkiewicz. 1993. [Building a Large Annotated Corpus of English: The Penn Treebank](https://aclanthology.org/J93-2004). *Computational Linguistics*, 19(2):313–330 (<https://aclanthology.org/J93-2004>)

Universal Dependencies (UD)

- Universal dependencies (UD) is a framework for morphosyntactic annotation of human language
 - used to create treebanks for more than 100 languages
 - facilitates multilingual parser development
 - cross-lingual learning
 - and parsing research from a language typology perspective
- Find out more at: <https://universaldependencies.org/>

Marie-Catherine de Marneffe, Christopher D. Manning, Joakim Nivre, Daniel Zeman; Universal Dependencies. *Computational Linguistics* 2021; 47 (2): 255–308. doi: https://doi.org/10.1162/coli_a_00402

UD CoNLL-U Format

- Annotations are encoded in plain text files (UTF-8) with three types of lines:
 - Word lines containing the annotation of a word/token in 10 fields separated by single tab characters; see below.
 - Blank lines marking sentence boundaries.
 - Comment lines starting with hash (#).
- More at: <https://universaldependencies.org/format.html>

Sabine Buchholz and Erwin Marsi. 2006. [CoNLL-X Shared Task on Multilingual Dependency Parsing](https://universaldependencies.org/format.html). In *Proceedings of the Tenth Conference on Computational Natural Language Learning (CoNLL-X)*, pages 149–164, New York City. Association for Computational Linguistics. <https://aclanthology.org/W06-2920>

Word lines in CoNLL-U Format

- Sentences consist of one or more word lines, and word lines contain the following fields:
 - ID: Word index, integer starting at 1 for each new sentence; may be a range for multiword tokens; may be a decimal number for empty nodes (decimal numbers can be lower than 1 but must be greater than 0). **We will ignore multiword and empty tokens!**
 - FORM: Word form or punctuation symbol.
 - LEMMA: Lemma or stem of word form.
 - UPOS: [Universal part-of-speech tag](#).
 - XPOS: Language-specific part-of-speech tag; underscore if not available.
 - FEATS: List of morphological features from the [universal feature inventory](#) or from a defined [language-specific extension](#); underscore if not available.
 - HEAD: Head of the current word, which is either a value of ID or zero (0).
 - DEPREL: [Universal dependency relation](#) to the HEAD ([root](#) iff HEAD = 0) or a defined language-specific subtype of one.
 - DEPS: Enhanced dependency graph in the form of a list of head-deprel pairs.
 - MISC: Any other annotation.

Example

```
Open [icon] en_partut-ud-train.conllu
~/Downloads


1 # sent_id = en_partut-ud-3
2 # text = Distribution of this license does not create an attorney-client relationship.
3 1 Distribution distribution NOUN S Number=Sing 7 nsubj _ _
4 2 of of ADP E _ 4 case _ _
5 3 this this DET DD Number=Sing|PronType=Dem 4 det _ _
6 4 license license NOUN S Number=Sing 1 nmod _ _
7 5 does do AUX VM Mood=Ind|Number=Sing|Person=3|Tense=Pres|VerbForm=Fin 7 aux _ _
8 6 not not PART PART Polarity=Neg 7 advmod _ _
9 7 create create VERB V Mood=Ind|Number=Plur|Tense=Pres|VerbForm=Fin 0 root _ _
10 8 an a DET RI Definite=Ind|Number=Sing|PronType=Art 12 det _ _
11 9 attorney attorney NOUN S Number=Sing 12 nmod _ SpaceAfter=No
12 10 - - PUNCT FF _ 9 punct _ SpaceAfter=No
13 11 client client NOUN S Number=Sing 9 compound _ _
14 12 relationship relationship NOUN S Number=Sing 7 obj _ SpaceAfter=No
15 13 . . PUNCT FS _ 7 punct _ _
16
17 # sent_id = en_partut-ud-4
18 # text = Creative Commons provides this information on an "as-is" basis.
19 1 Creative Creative PROPN SP _ 3 nsubj _ _
20 2 Commons Commons PROPN SP _ 1 flat _ _
21 3 provides provide VERB V Mood=Ind|Number=Sing|Person=3|Tense=Pres|VerbForm=Fin 0 root _ _
22 4 this this DET DD Number=Sing|PronType=Dem 5 det _ _
23 5 information information NOUN S Number=Sing 3 obj _ _
24 6 on on ADP E _ 13 case _ _
25 7 an a DET RI Definite=Ind|Number=Sing|PronType=Art 13 det _ _
26 8 " " PUNCT FB _ 11 punct _ SpaceAfter=No
27 9 as as ADP E _ 11 mark _ SpaceAfter=No
28 10 - - PUNCT FF _ 9 punct _ SpaceAfter=No
29 11 is be VERB V Mood=Ind|Number=Sing|Person=3|Tense=Pres|VerbForm=Fin 13 amod _ SpaceAfter=No
30 12 " " PUNCT FB _ 11 punct _ _
31 13 basis basis NOUN S Gender=Masc|Number=Sing 3 obl _ SpaceAfter=No
32 14 . . PUNCT FS _ 3 punct _ _
33
34 # sent_id = en_partut-ud-5
35 # text = Creative Commons makes no warranties regarding the information provided, and disclaims liability for damages resulting from its use.
36 1 Creative Creative PROPN SP _ 3 nsubj _ _
37 2 Commons Commons PROPN SP _ 1 flat _ _
38 3 makes make VERB V Mood=Ind|Number=Sing|Person=3|Tense=Pres|VerbForm=Fin 0 root _ _
39 4 no no DET DI PronType=Ind 5 det _ _
40 5 warranties warranty NOUN S Number=Plur 3 obj _ _
41 6 regarding regard VERB V Number=Sing|Tense=Pres|VerbForm=Part 5 acl _ _
42 7 the the DET RD Definite=Def|PronType=Art 8 det _ _
43 8 information information NOUN S Number=Sing 6 obj _ _
44 9 provided provide VERB V Tense=Past|VerbForm=Part 8 acl _ SpaceAfter=No
45 10 , , PUNCT FF _ 12 punct _ _
46 11 and and CCONJ CC _ 12 cc _ _
```

UPOS

Universal part-of-speech tags (UPOS). Typos and abbreviations are given the category of the unabbreviated or correct word.

Traditional POS	UPOS	Category
noun	NOUN	common noun
	PROPN	proper noun
verb	VERB	main verb
	AUX	auxiliary verb or other tense, aspect, or mood particle
adjective	ADJ	adjective
	DET	determiner (including article)
	NUM	numeral (cardinal)
adverb	ADV	adverb
pronoun	PRON	pronoun
preposition	ADP	adposition (preposition/postposition)
conjunction	CCONJ	coordinating conjunction
	SCONJ	subordinating conjunction
interjection	INTJ	interjection
–	PART	particle (special single word markers in some languages)
–	X	other (e.g., words in foreign language expressions)
–	SYM	non-punctuation symbol (e.g., a hash (#) or emoji)
–	PUNCT	punctuation

Example of a tree for UD in CoNLL-U format

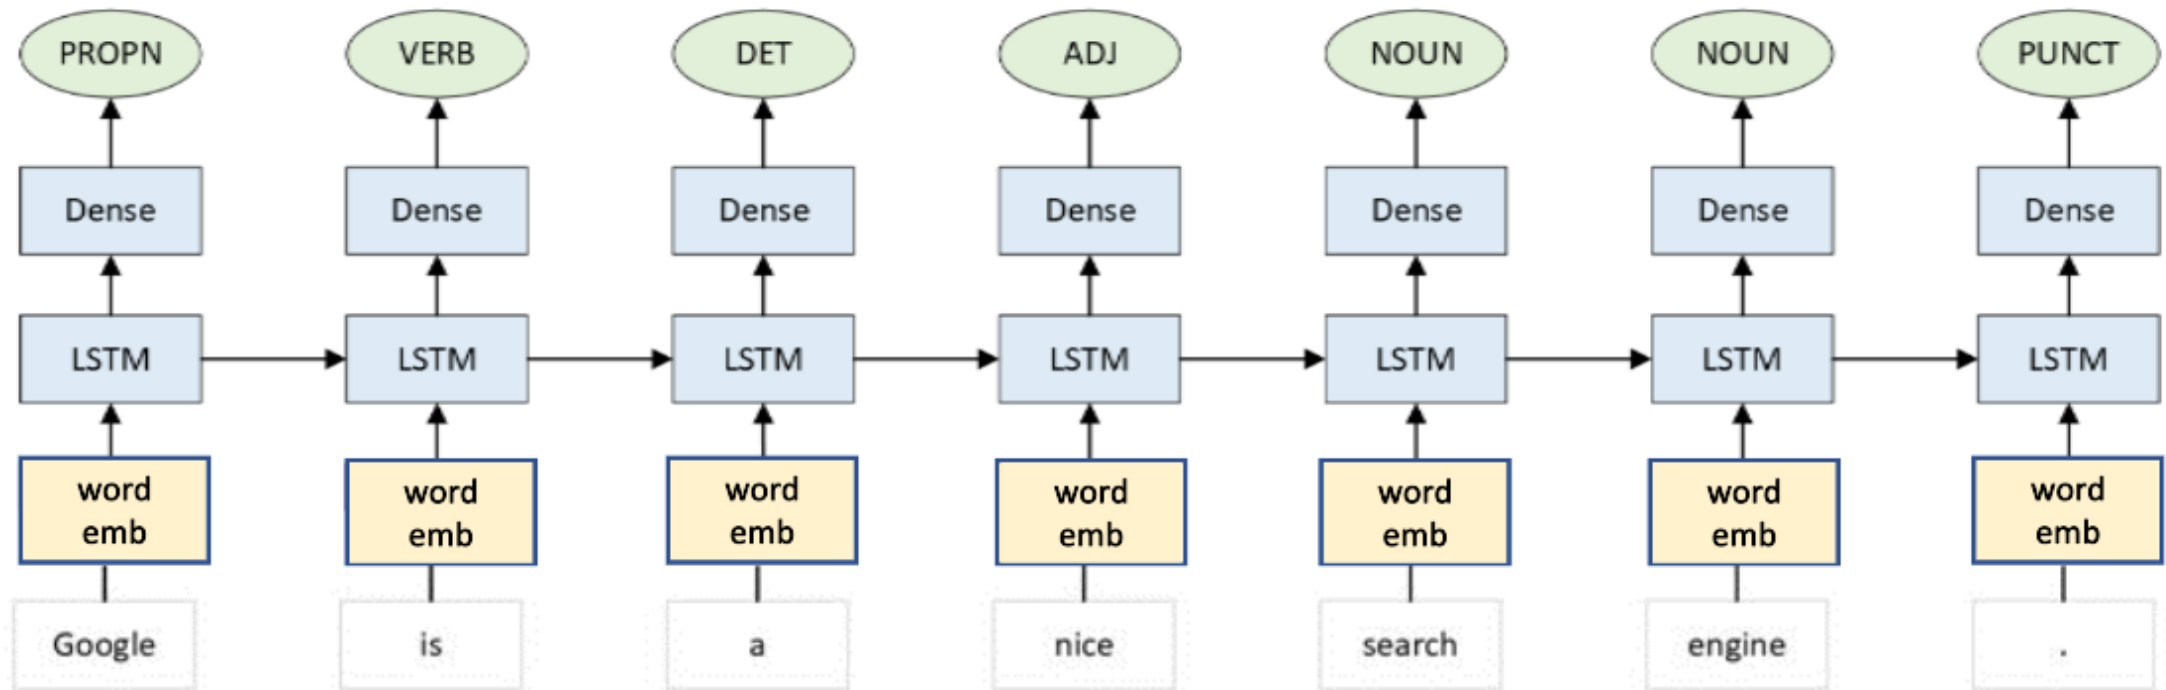


1	Cathy	Cathy	N	N	eigen ev neut	2	su7
2	zag	zie	V	V	trans ovt 1of2of3 ev	0	ROOT
3	hen	hen	Pron	Pron	per 3 mv datofacc	2	obj1
4	wild	wild	Adj	Adj	attr stell onverv	5	mod
5	zwaaien	zwaai	N	N	soort mv neut	2	vc
6	.	.	Punc	Punc	punct	5	punct

UD Treebanks datasets

- English: https://github.com/UniversalDependencies/UD_English-EWT/tree/master
- Galician: https://github.com/UniversalDependencies/UD_Galician-TreeGal/tree/master
- And many more: <https://universaldependencies.org/#download>

Architecture



Tips and hints

Organize your code

- Include a user manual explaining how to run
- Define classes to create reusable pieces of code
 - See our tutorial "Python Basics introduction and OOP" available in the Virtual Campus.
- Your main notebook should make easy to relaunch all the relevant tasks: training, evaluation and generation of labels, etc.
- You may have several .py files that can be imported in your main notebook (see next slide)
- Consider saving and loading your good trained models.
- Include a brief discussion of the implementation decisions, differences across the evaluated models that you might have explored, as well as an analysis of the performance across the evaluated datasets
 - In a separate PDF not exceeding 3 pages and written using Calibri font style and a size of minimum 11pt. Put your names on it.
 - You may include the necessary comments in python comments/jupyter notebook text cells

Example: using .py files in your notebook

The screenshot displays a Jupyter Notebook interface with the title "myP1notebook.ipynb". The top menu bar includes "File", "Edit", "View", "Insert", "Runtime", "Tools", and "Help", with a status "All changes saved". On the right, there are buttons for "Comment", "Share", and a settings icon. Below the menu, a toolbar shows "RAM" and "Disk" usage, and a status "Editing".

The left sidebar, titled "Files", shows a file explorer with a tree structure: "drive" (expanded) containing "MyDrive" (expanded) and "Ap1", which contains the file "mymodule.py".

The main notebook area contains the following content:

We can import a .py file in the usual way. We just need to be careful with paths, adjusting them as necessary. You should mount and unmount your google drive folder by using the button in the toolbar on the left.

```
from drive.MyDrive.Ap1.mymodule import MyTagger
```

[2] mytaggger = MyTagger()

Notice that it could be better to have a single folder where we store everything: the notebook, dataset files and all .py module files.

A convenient way to proceed is to change the PATH to include such folder:

```
# add a specific gdrive folder to the PATH
import sys
sys.path.append('/content/gdrive/mypythondirectory')
```

and then to import like this:

```
from mymodule import MyTagger
```

[3] import sys
sys.path.append('/content/drive/MyDrive/Ap1/')

[4] from mymodule import MyTagger
myTagger = MyTagger()

On the right, a separate editor window titled "mymodule.py" shows the code for the module:

```
1 import keras
2 import tensorflow as tf
3 import numpy as np
4 import tqdm
5
6 class MyTagger(object):
7
8     def build_model(self):
9         pass
10
11     def train(self, train_sents, dev_sents):
12         pass
13
14     def evaluate(self, eval_sents):
15         pass
16
17     def predict(self, test_sents):
18         pass
19
```

Example: using .py files in your notebook

The screenshot displays a Jupyter Notebook interface with the title "myP1notebook.ipynb". The top menu bar includes "File", "Edit", "View", "Insert", "Runtime", "Tools", and "Help", with a status "All changes saved". On the right, there are buttons for "Comment", "Share", and a settings icon. Below the menu, a toolbar shows "RAM" and "Disk" usage, and an "Editing" mode indicator.

The left sidebar, titled "Files", shows a file explorer with a tree view containing "drive", "MyDrive", and "Ap1" folders, with a file "mymodule.py" located inside "Ap1".

The main notebook area contains the following content:

We can import a .py file in the usual way. We just need to be careful with paths, adjusting them as necessary. You should mount and unmount your google drive folder by using the button in the toolbar on the left.

```
from drive.MyDrive.Ap1.mymodule import MyTagger
```

[2] mytaggger = MyTagger()

Notice that it could be better to have a single folder where we store everything: the notebook, dataset files and all .py module files.

A convenient way to proceed is to change the PATH to include such folder:

```
# add a specific gdrive folder to the PATH
import sys
sys.path.append('/content/gdrive/mypythondirectory')
```

and then to import like this:

```
from mymodule import MyTagger
```

[3] import sys
sys.path.append('/content/drive/MyDrive/Ap1/')

[4] from mymodule import MyTagger
myTagger = MyTagger()


On the right, a separate editor window titled "mymodule.py" shows the code for the module:

```
1 import keras
2 import tensorflow as tf
3 import numpy as np
4 import tqdm
5
6 class MyTagger(object):
7
8     def build_model(self):
9         pass
10
11     def train(self, train_sents, dev_sents):
12         pass
13
14     def evaluate(self, eval_sents):
15         pass
16
17     def predict(self, test_sents):
18         pass
19
```


Recurrent neurons and layers




A recurrent neuron



Recurrent neuron unrolled through time




A layer of recurrent neurons




A layer of recurrent neurons unrolled through time

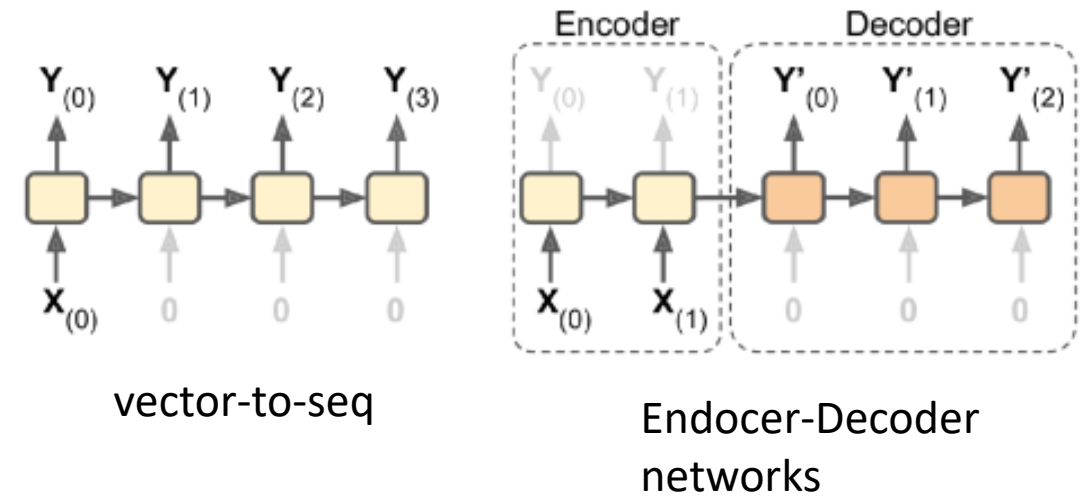

Recurrent neurons and layers




Like memory cells



h : hidden state




Deep RNNs



```
model = keras.models.Sequential([
    keras.layers.SimpleRNN(20, return_sequences=True, input_shape=[None,
1]),
    keras.layers.SimpleRNN(20, return_sequences=True),
    keras.layers.SimpleRNN(1)
])
```

```
model = keras.models.Sequential([
    keras.layers.SimpleRNN(20, return_sequences=True, input_shape=[None,
1]),
    keras.layers.SimpleRNN(20),
    keras.layers.Dense(1)
])
```

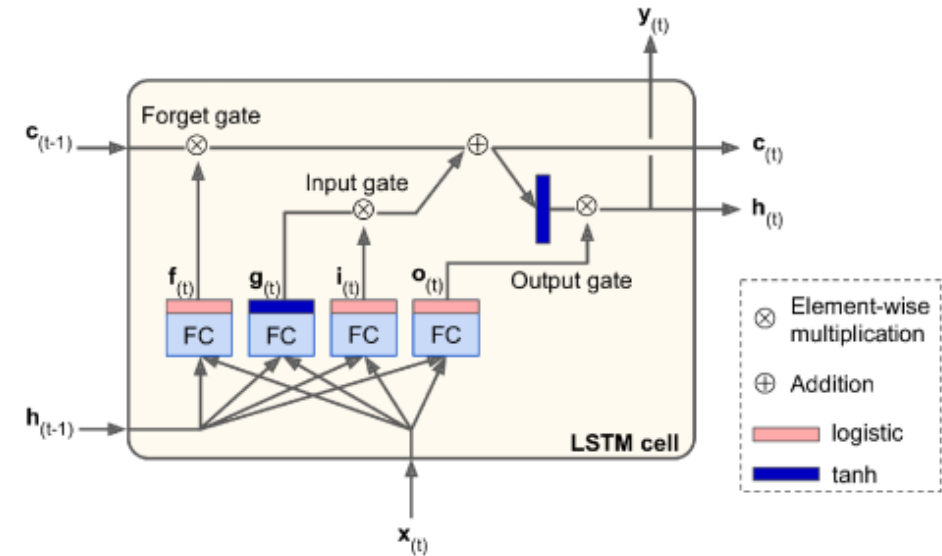
Bidirectional RNNs



LSTM cells

LSTM as a black box:

- used like any other basic memory cell
- much better performance and faster convergence
- detecting long-term dependencies in the data (successful at capturing long-term patterns)

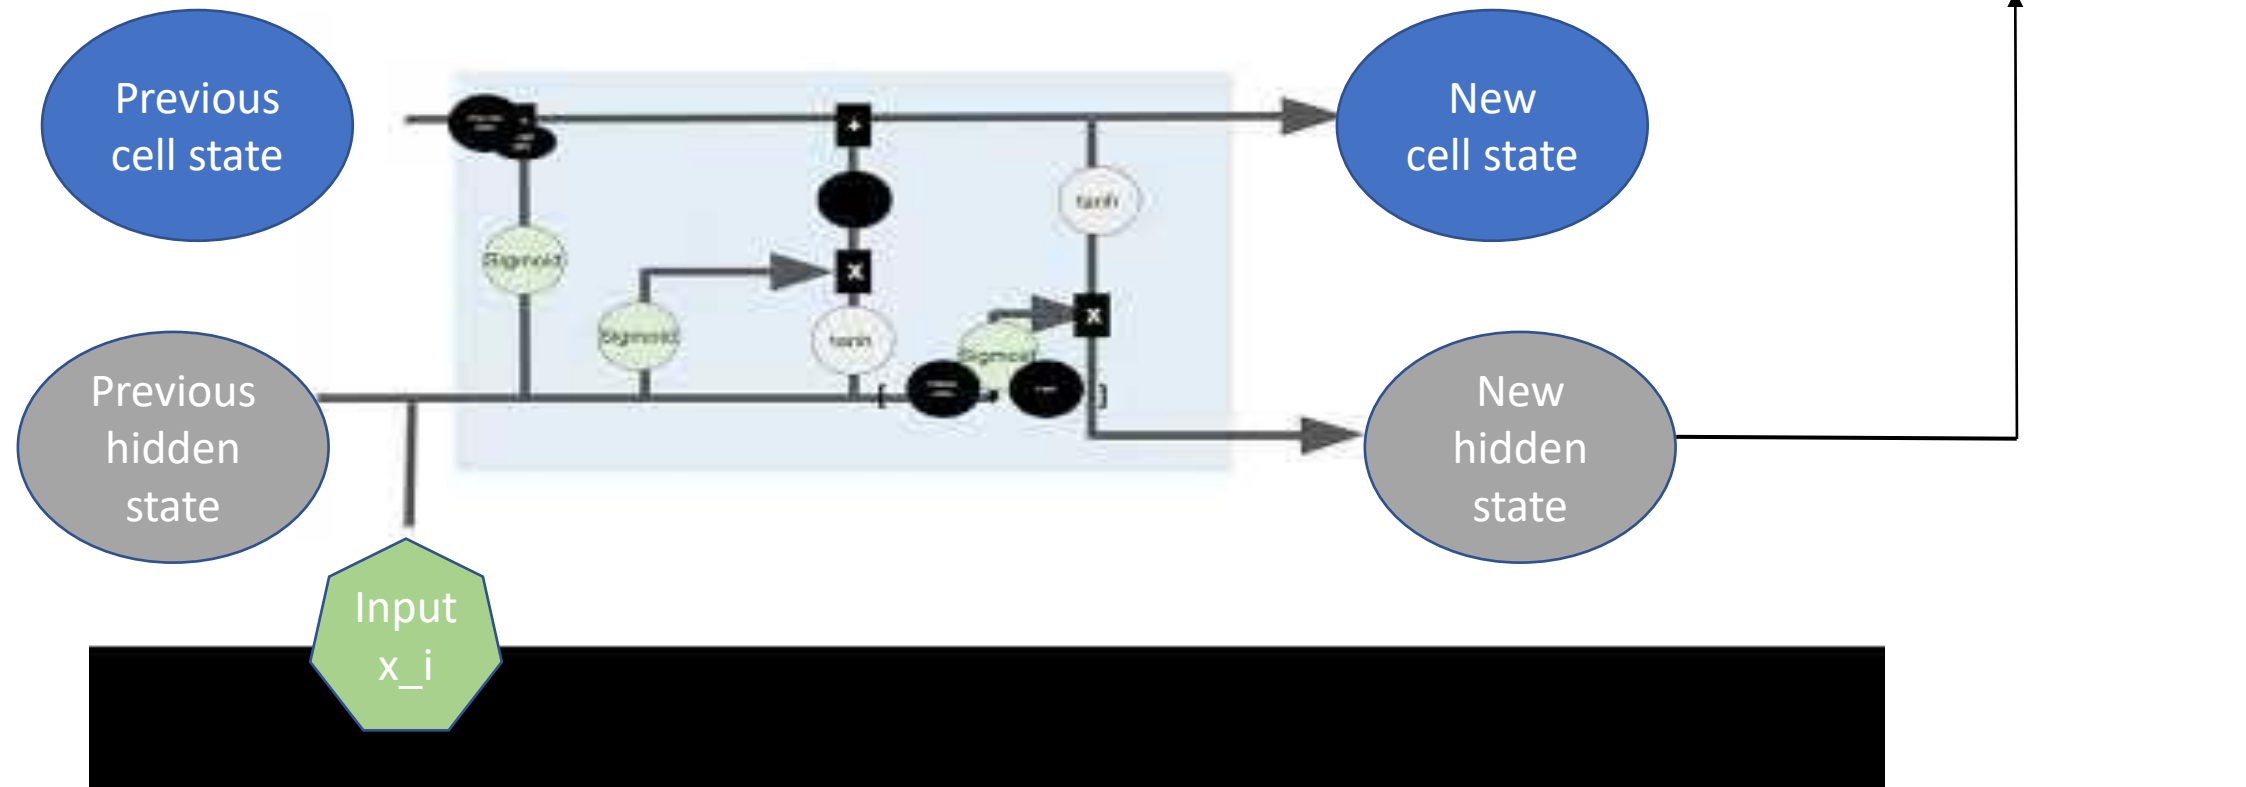


```
model = keras.models.Sequential([
    keras.layers.SimpleRNN(20, return_sequences=True, input_shape=[None,
1]),
    keras.layers.SimpleRNN(20, return_sequences=True),
    keras.layers.TimeDistributed(keras.layers.Dense(10))
])
```



```
model = keras.models.Sequential([
    keras.layers.LSTM(20, return_sequences=True, input_shape=[None, 1]),
    keras.layers.LSTM(20, return_sequences=True),
    keras.layers.TimeDistributed(keras.layers.Dense(10))
])
```


LSTM cells



Defining the neural model in Keras

Creating the input samples

Preprocessing the treebank: Ignore sentences longer than 128 from the training, development and test sets. This means an input to the model will have the shape (max_sentence_length, embedding_size).

Recommendations:

1. Use the Tokenizer or TextVectorizer layer approaches.
2. Implement your models using the functional API.
3. Make sure an extra ID is used for unknown tokens (read the docs for the Tokenizer or the TextVectorizer layer)

Main steps to train and evaluate the word-level model:

1. **Transform the input sentences (strings) to a list of numerical IDs.**
2. **Pad the input sentences**, so they all have the same length (required for Keras models).
3. **Convert the output labels to IDs as well**, so we can train the model with the fit function.

Creating the input samples

texts_to_sequences (with our
trained Tokenizer), TextVectorizer
or an ad-hoc function

Mary has a cat

PROPN VERB DET NOUN

17 56 28 100

1 3 6 2

Spiderman exists

PROPN VERB

7 324

1 3

I am real

PRONOUN VERB ADJ

98 76 3

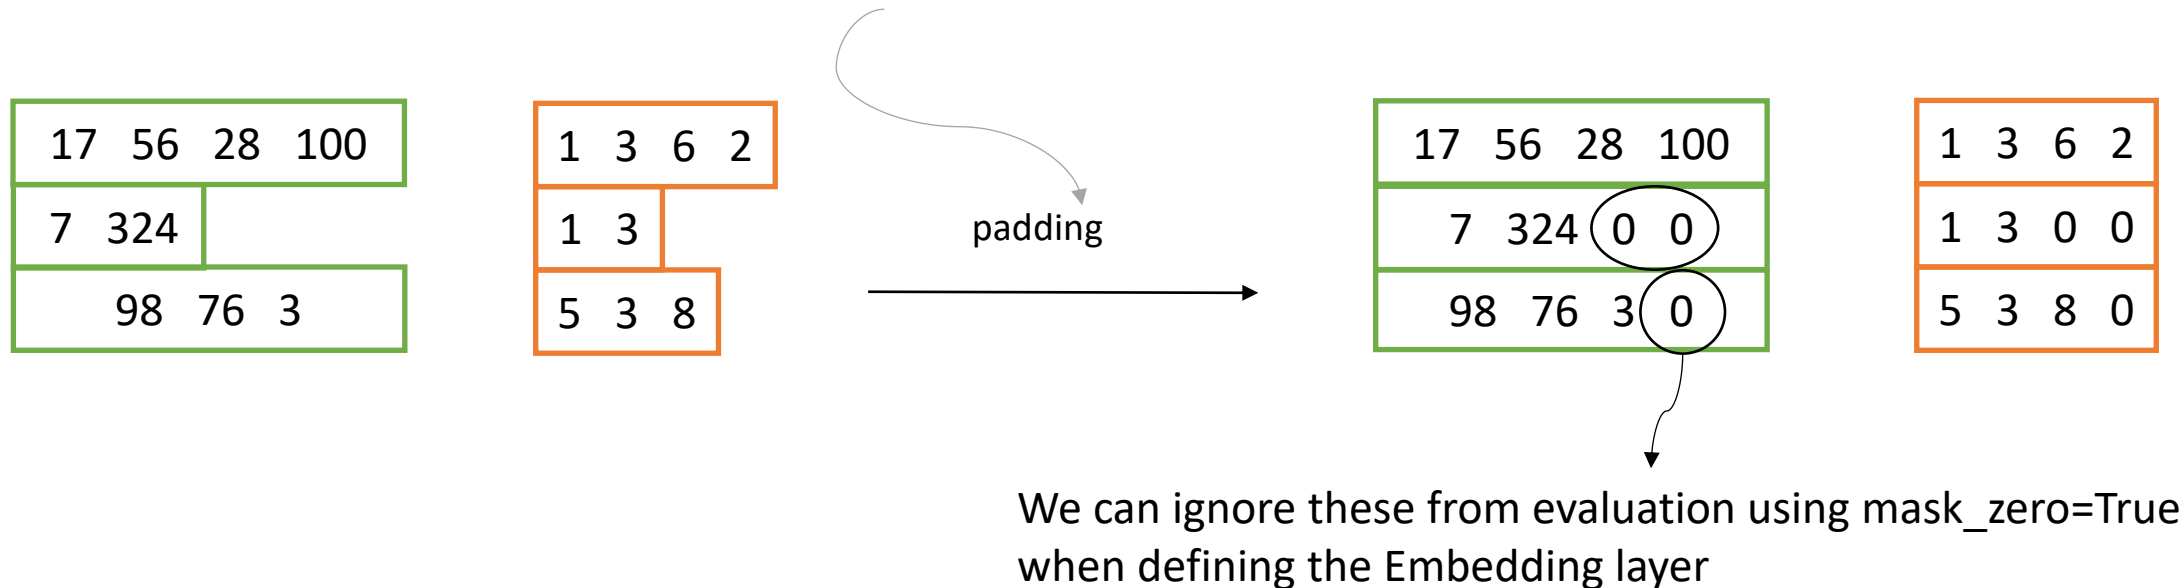
5 3 8

Creating the input samples

In a single batch, we have input sentences with different lengths.

All sentences in a batch must have the same length in order to be fed to the keras models.

The solution is to apply padding to the inputs




Creating the input samples

We are addressing a multi-label classification problem.

We will be optimizing our model with a categorical_crossentropy loss:

https://www.tensorflow.org/api_docs/python/tf/keras/metrics/categorical_crossentropy

To do so, we need to transform the desired output label (initially represented as a string) into a one-hot vector to encode which class the model must predict.



Alternatively to a categorical_crossentropy loss we also can use a sparse categorical crossentropy loss (in such case, we do not need to use the to_categorical function):

https://www.tensorflow.org/api_docs/python/tf/keras/losses/SparseCategoricalCrossentropy

Declaring the model – Relevant layers

Input layer: https://www.tensorflow.org/api_docs/python/tf/keras/Input [seen in P0]

Embedding layer: https://www.tensorflow.org/api_docs/python/tf/keras/layers/Embedding [seen in P0]

LSTM layer: https://www.tensorflow.org/api_docs/python/tf/keras/layers/LSTM [seen in P0]

Bidirectional layer: https://www.tensorflow.org/api_docs/python/tf/keras/layers/Bidirectional [optional]

Dense layer: https://www.tensorflow.org/api_docs/python/tf/keras/layers/Dense [seen in P0]

TimeDistributed layer: https://www.tensorflow.org/api_docs/python/tf/keras/layers/TimeDistributed [not seen in P0, but relevant for P1]

tf.keras.layers.TimeDistributed

A wrapper to apply a layer to every temporal slice of an input.

The dimension of index one of the input is considered the temporal dimension.

In our case, the shape of the input for the word-level PoS tagger is (batch_size, max_sentence_length, word_embedding_size).

Given a TimeDistributed layer, it will apply the wrapped layer to every element of the sentence.

tf.keras.layers.TimeDistributed

model = ...

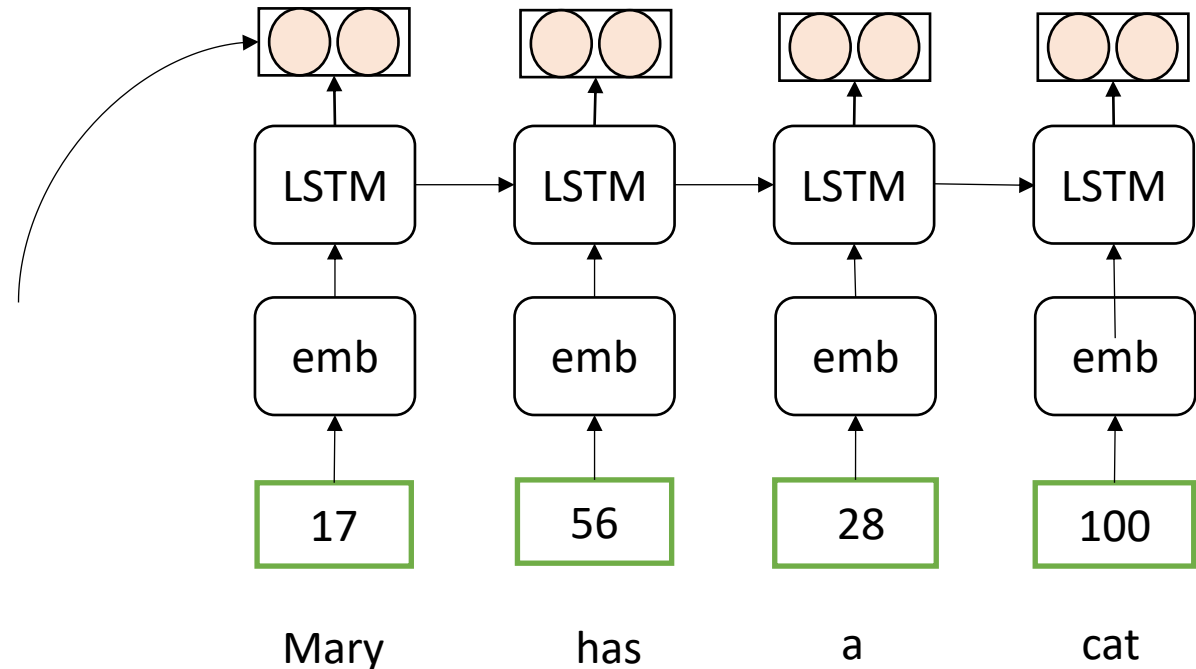
...

We add the embedding layer

We add the LSTM layer

...

The output of the LSTM must be a vector for each word.



tf.keras.layers.TimeDistributed

model = ...

...


We add the embedding layer

We add the LSTM layer

...

The output of the LSTM must be a vector for each word, i.e. the hyperparameter `return_sequences=True` for the LSTM instance

However, the output layer (a Dense layer) is not intended to be applied over sequences, how can we deal with that? Using the `TimeDistributed` wrapper, which does exactly that.



tf.keras.layers.TimeDistributed

model = ...

...

We add the embedding layer

We add the LSTM layer


...

```
model.add(TimeDistributed(Dense(nlabels,  
activation='softmax')))
```

The output of the LSTM must be a vector for each word.

...

model.fit(...)



tf.keras.layers.TimeDistributed

model = ...

...

We add the embedding layer

We add the LSTM layer

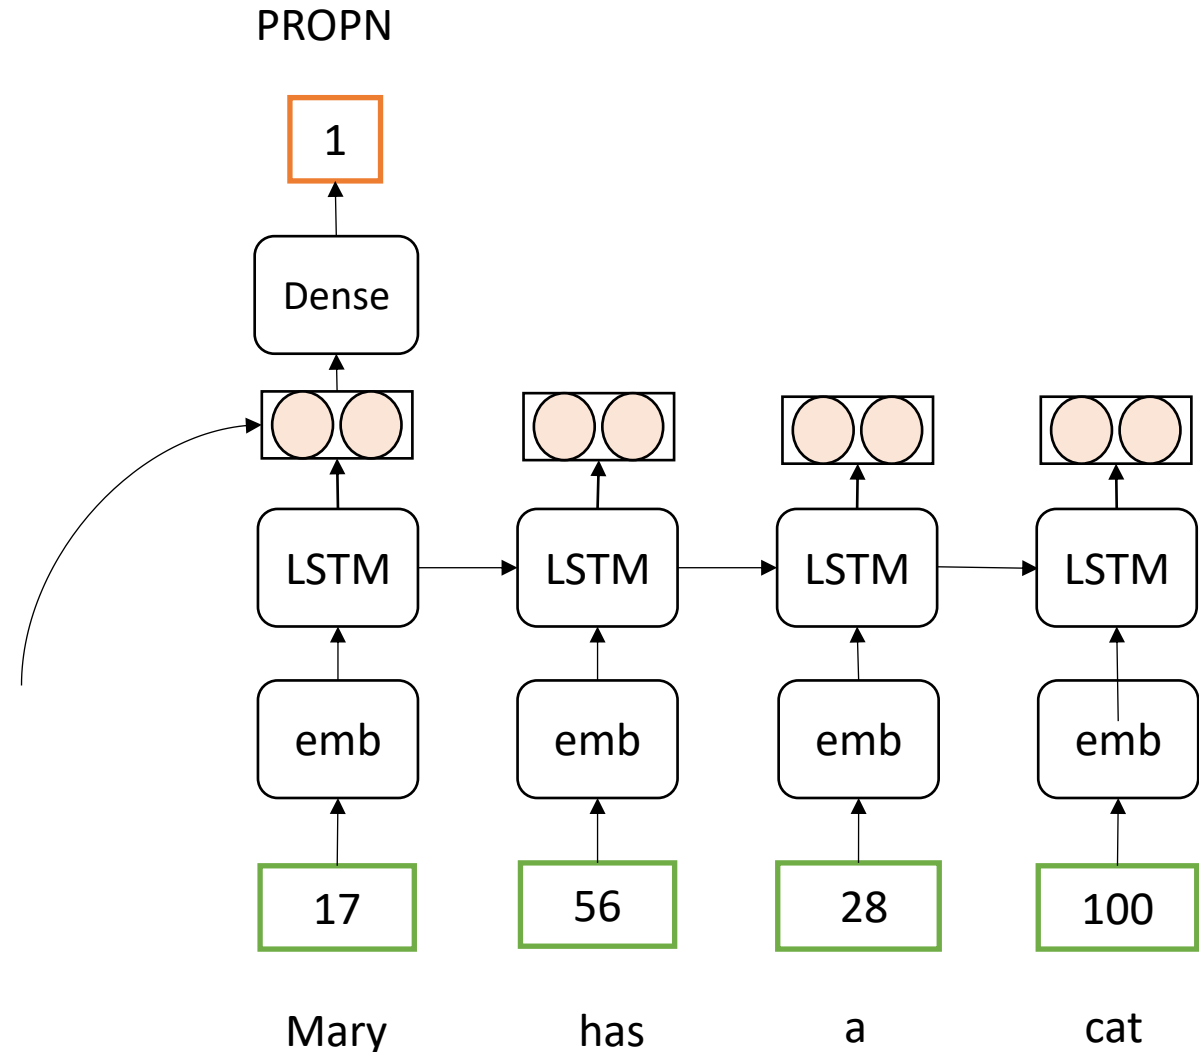
...

```
model.add(TimeDistributed(Dense(nlabels,  
activation='softmax')))
```

The output of the LSTM must be a vector for each word.

...

model.fit(...)



tf.keras.layers.TimeDistributed

model = ...

...

We add the embedding layer

We add the LSTM layer

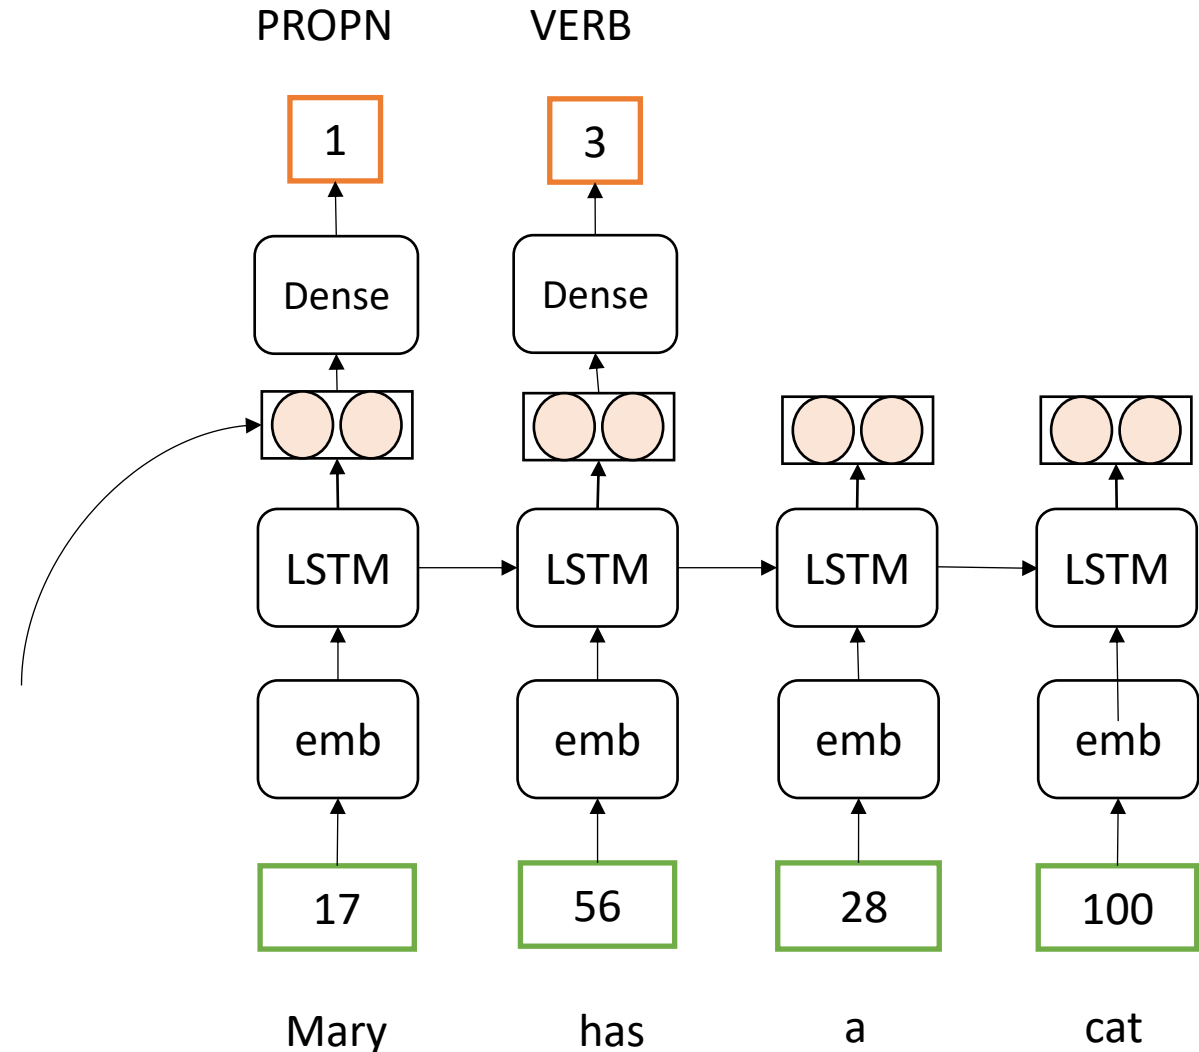
...

```
model.add(TimeDistributed(Dense(nlabels,  
activation='softmax')))
```

The output of the LSTM must be a vector for each word.

...

model.fit(...)



tf.keras.layers.TimeDistributed

model = ...

...

We add the embedding layer

We add the LSTM layer

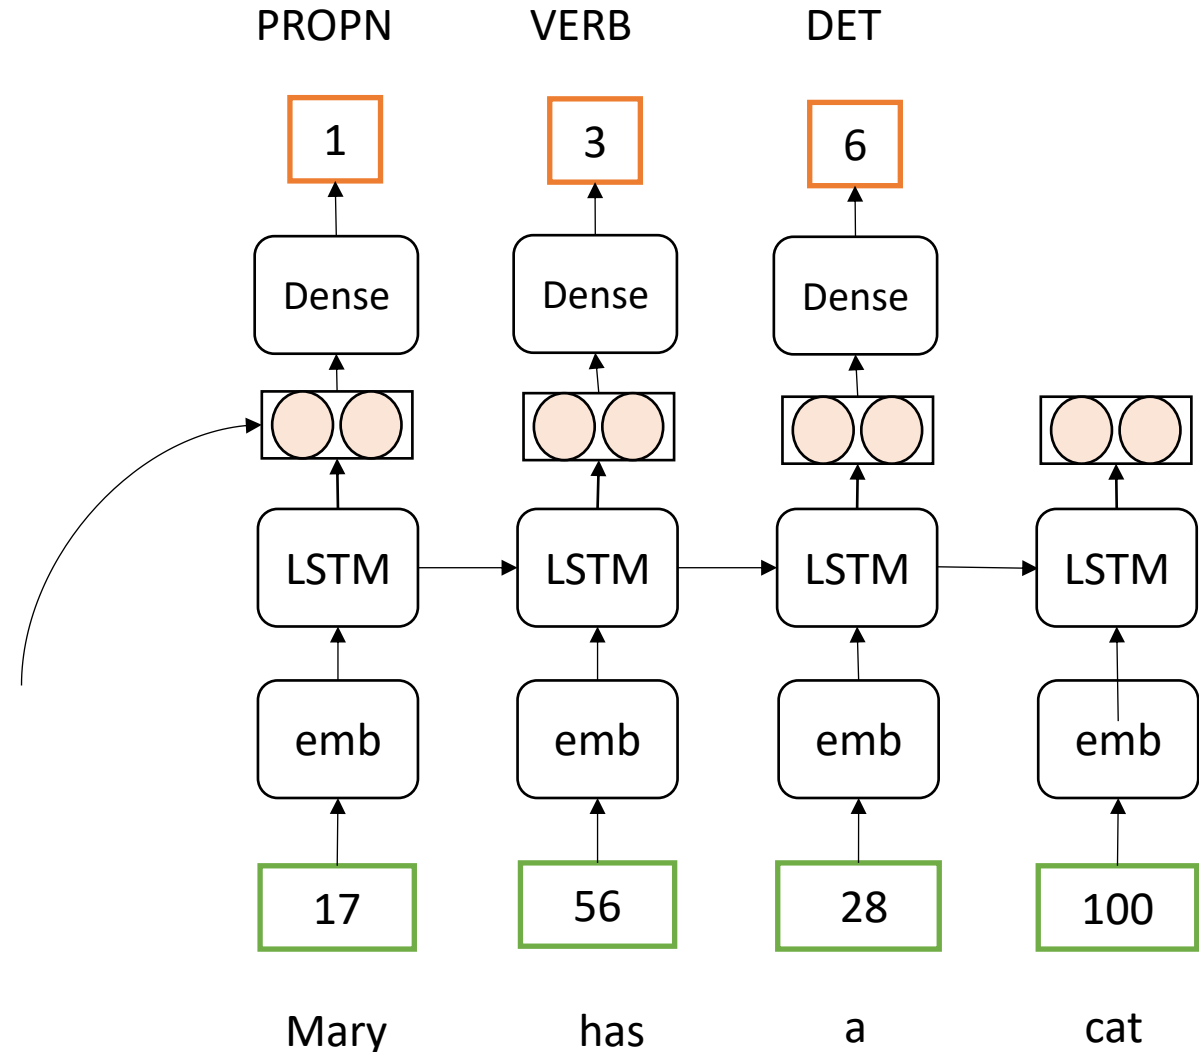
...

```
model.add(TimeDistributed(Dense(nlabels,  
activation='softmax')))
```

The output of the LSTM must be a vector for each word.

...

model.fit(...)



tf.keras.layers.TimeDistributed

model = ...

...

We add the embedding layer

We add the LSTM layer

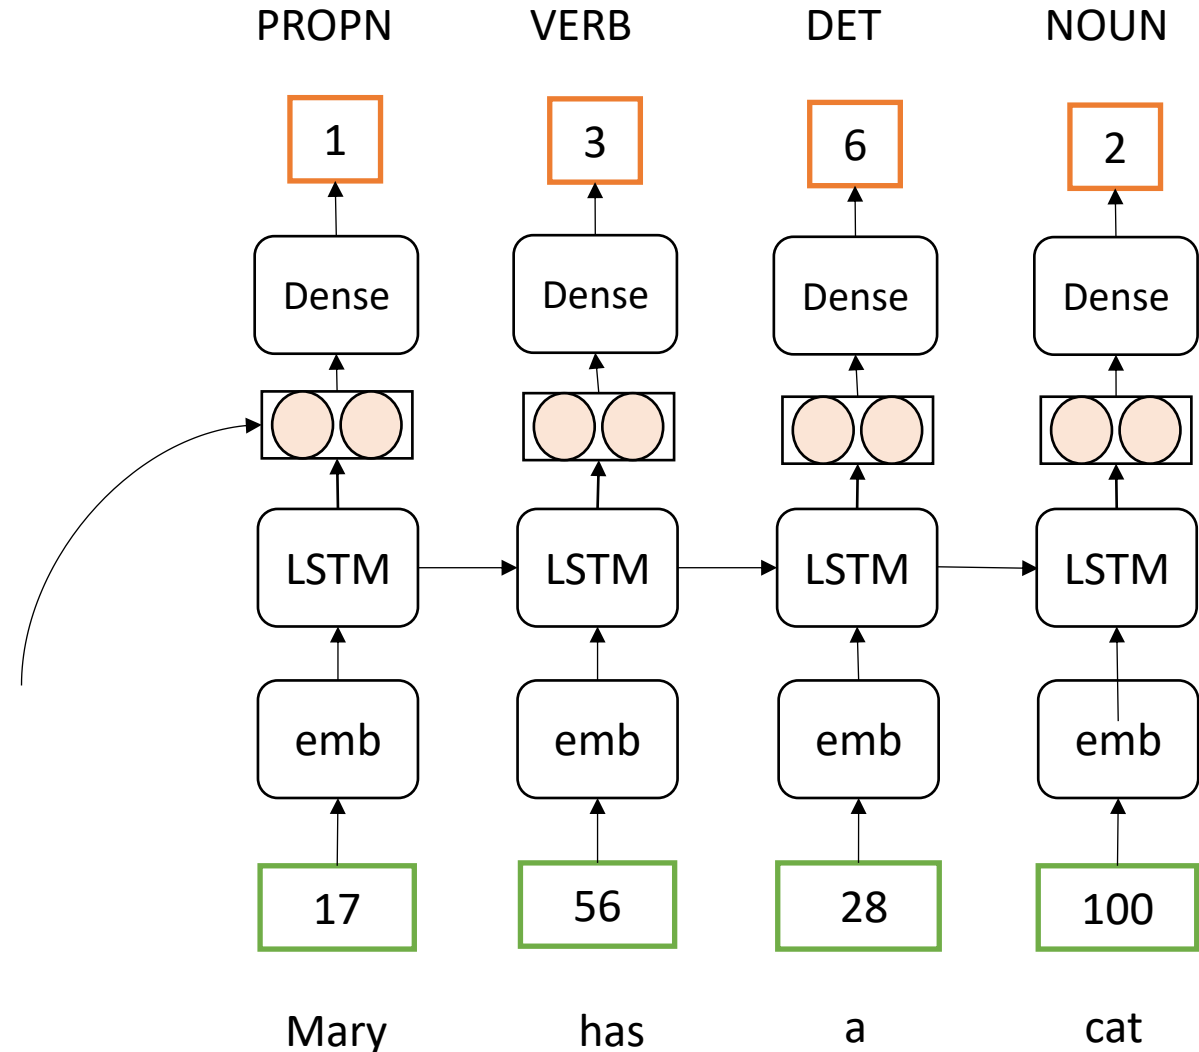
...

```
model.add(TimeDistributed(Dense(nlabels,  
activation='softmax')))
```

The output of the LSTM must be a vector for each word.

...

model.fit(...)



Proposed planning

Week	Minimum amount of tasks that you should have completed at such point
September 29- October 3	Read the assignment and check materials (starts on October 1)
October 6 – 10	Implement functions to process the UD English treebank, remove multiword and empty lines.
October 13 – 17	Implement functions for tokenization and mapping samples to ids
October 20 – 24	Implement functions for training and evaluation of the model. Consider some variants of the network modifying the parameters.
October 27 – 31	Implement generation of labels for unseen sentences. Test other non-English treebanks.
Until November 3 (Monday, 15.30!!)	Writing the report and submission